Mini project using Snowflake.

Title: Library Management System Analytic

Objective: To analyze and visualize data from a library management system using Snowflake Data Platform, incorporating multiple fact and dimension tables for comprehensive insights.

Schema: Star schema with multiple fact and dimension tables:
Fact Tables: This table stores quantitative data like book checkouts and monetary transactions related to fines associated with checkouts.
checkoutsfact (checkoutid, bookid, borrowerid, checkoutdate, returndate)
finepaymentsfact (paymentid, checkoutid, paymentdate, amount)

Dimension Tables: Contain descriptive attributes.
bookdim (bookid, title, author, genre, publishyear)
borrowerdim (bor_index, borrowerid, name, address)
librarybranchdim (branchid, branchname, location)
datedim (dateid, day, month, year, weekday)

The checkoutsfact table has relationships with the bookdim, borrowerdim, librarybranchdim, and datedim tables. It links to bookdim through bookid, to borrowerdim through borrowerid, to librarybranchdim through branchid, and to datedim through checkoutdate.

The finepaymentsfact table has a relationship with the checkoutsfact table through checkoutid, and it indirectly links to other dimension tables through its relationship with checkoutsfact. Each dimension table (bookdim, borrowerdim, librarybranchdim, datedim) serves as a hub for details related to books, borrowers, library branches, and dates, respectively.

Process:

1. Logged into snowflake and created new database. Used below command to use the new database,
   USE DATABASE library_management;

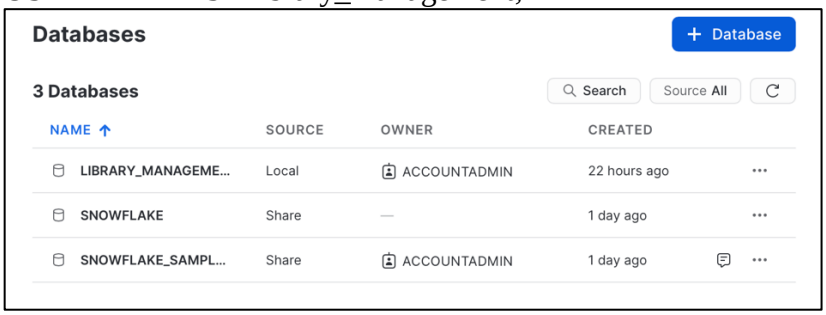

2. Executed CREATE statement in Snowflake SQL worksheet:

```
CREATE TABLE bookdim (
    book_id VARCHAR(255) PRIMARY KEY,
    book_title VARCHAR(255),
```

```
    book_author VARCHAR(255),
    year_of_publication INT,
    publisher VARCHAR(225)
);


CREATE TABLE borrowerdim1 (
    bor_index INT PRIMARY KEY,
    borrower_name VARCHAR(255),
    borrower_area VARCHAR(255),
    borrower_id INT
);

CREATE TABLE librarybranchdim (
    branch_id INT PRIMARY KEY,
    branchname VARCHAR(255),
    branch_address VARCHAR(255),
    branch_city VARCHAR(255),
    branch_phone VARCHAR(255),
    branch_zip INT
);

CREATE TABLE datedim (
    dateid INT PRIMARY KEY,
    day INT,
    month INT,
    year INT,
    weekday VARCHAR(20)
);
```

Snowflake stores DATE data more efficiently than VARCHAR, providing better query performance. Hence, I have used DATE datatype with format (ISO YYYY-MM-DD)

```
CREATE TABLE checkoutsfact (
    checkout_id INT PRIMARY KEY,
    book_id VARCHAR(255),
    borrower_id INT,
    checkoutdate DATE,
    returndate DATE,
    FOREIGN KEY (book_id) REFERENCES bookdim(book_id),
    FOREIGN KEY (borrower_id) REFERENCES borrowerdim(borrower_id)
);


CREATE TABLE finepaymentsfact (
    paymentid INT PRIMARY KEY,
```

checkoutid INT,
paymentdate DATE,
amount DECIMAL(10, 2),
FOREIGN KEY (checkoutid) REFERENCES checkoutsfact(checkoutid)
);

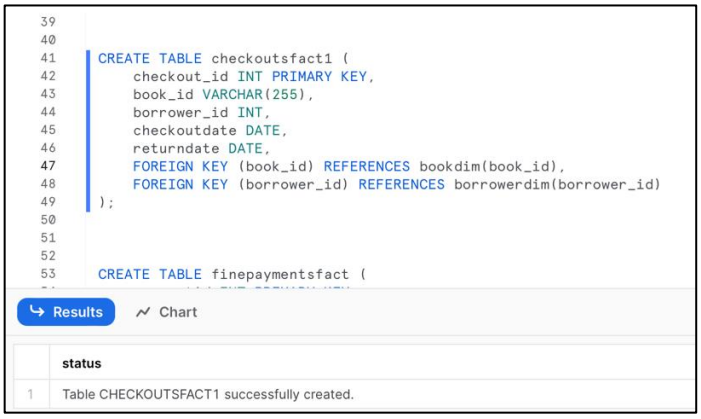

3. Loaded data from csv to tables in snowflake.

- Created a stage using below command,
  CREATE STAGE library;
- Files are first copied (“staged”) to an internal (Snowflake) stage, then loaded into a table.
- Uploaded CSV file from local machine to snowflake using worksheet query-
  PUT *file:///Users/vaidehipatel/Documents/Course\ Sem\ 1/DBMS/Snowflake/books.csv @library;*
- Executed COPY query to load data from csv to tables in snowflake.

Leveraged Snowflake's COPY command for efficient bulk loading.

CREATE FILE FORMAT csv_format
TYPE = CSV
FIELD_OPTIONALLY_ENCLOSED_BY = ""
NULL_IF = ('');

```
COPY INTO bookdim
FROM @library/books.csv
FILE_FORMAT = (FORMAT_NAME = 'csv_format'
FIELD_OPTIONALLY_ENCLOSED_BY = '"');

COPY INTO borrowerdim
FROM @library/borrower_data.csv
FILE_FORMAT = (FORMAT_NAME = 'csv_format'
FIELD_OPTIONALLY_ENCLOSED_BY = '"');

COPY INTO librarybranchdim
FROM @library/queens_library_branches.csv
FILE_FORMAT = (FORMAT_NAME = 'csv_format'
FIELD_OPTIONALLY_ENCLOSED_BY = '"');

COPY INTO checkoutsfact
FROM @library/checkouts_data.csv
FILE_FORMAT = (FORMAT_NAME = 'csv_format' DATE_FORMAT = 'YYYY-MM-DD');
```

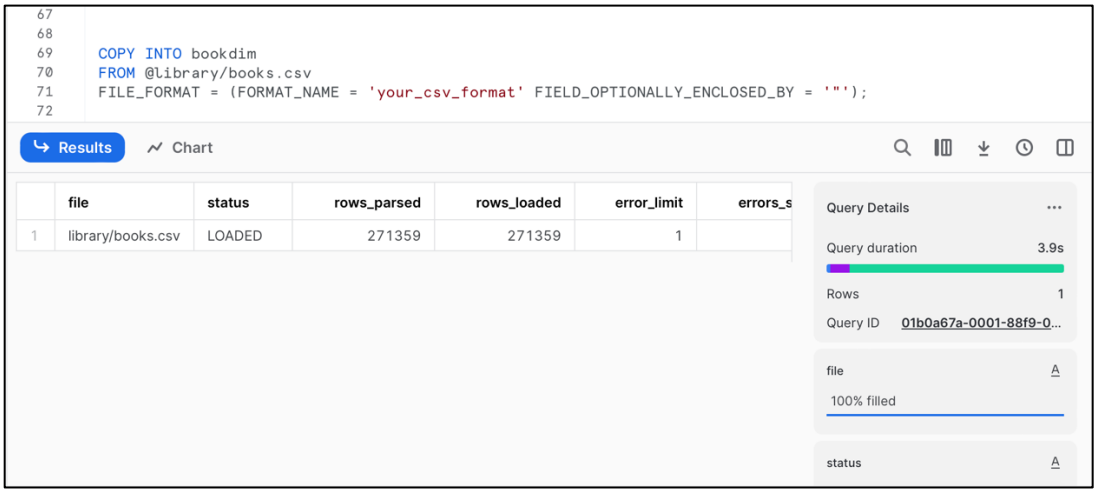

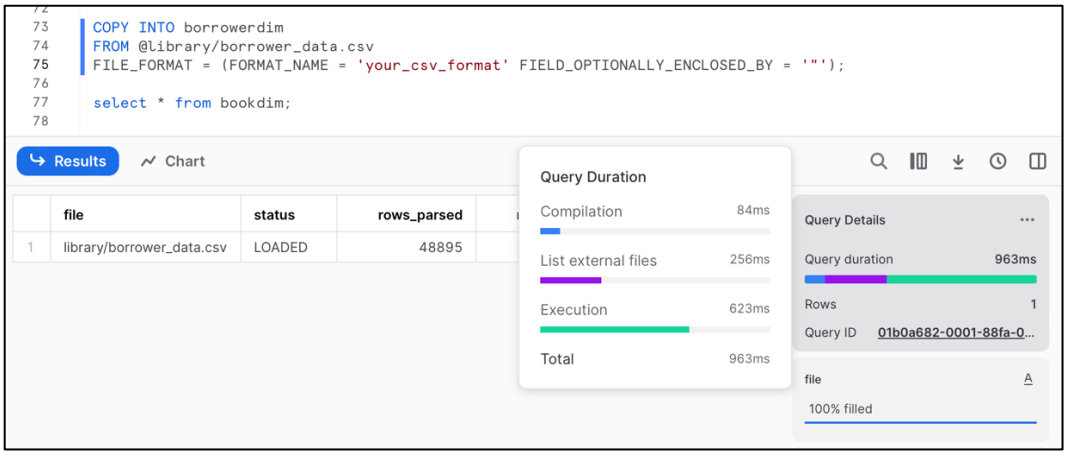

Loaded other tables using ‘Load a file’ button from Stage into an existing table using Snowsight.

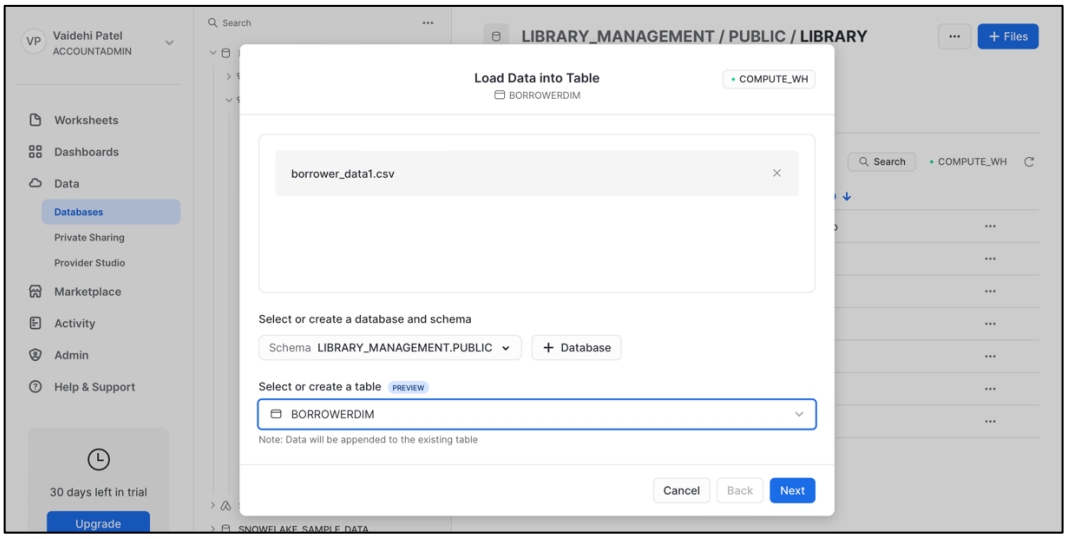

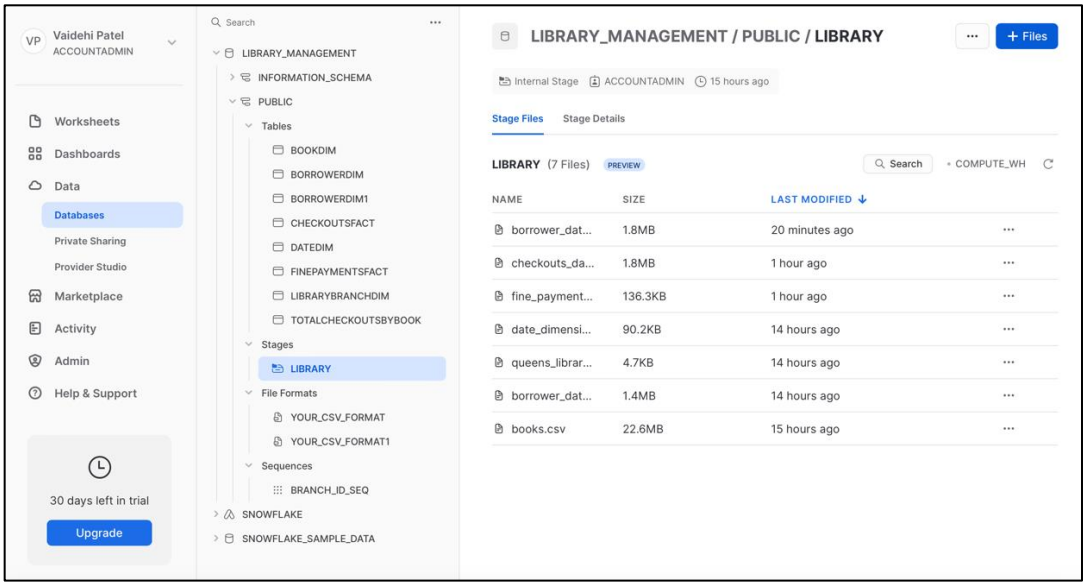

4. **Clustering**

Leveraged Snowflake's automatic clustering to organize the data based on book_id key. This improves query performance by physically storing related data together.

ALTER TABLE CheckoutsFact CLUSTER BY (book_id);

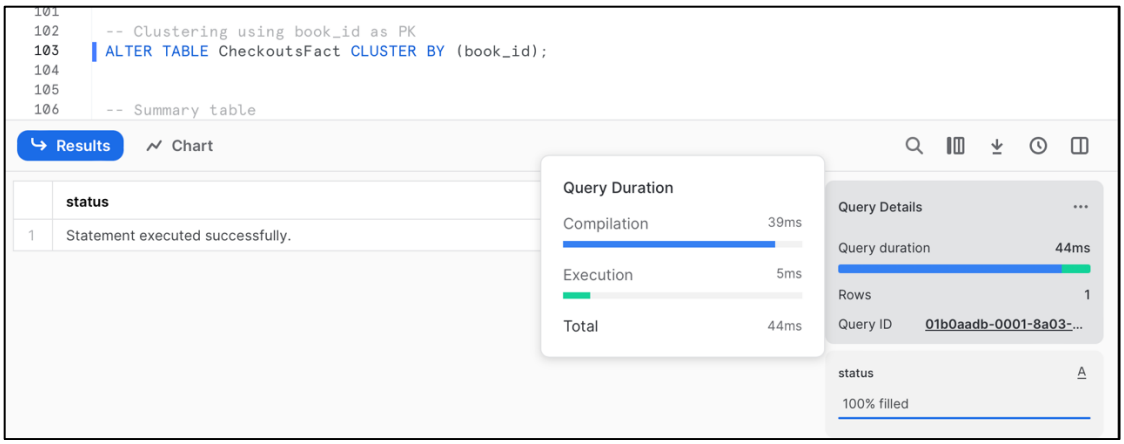

5. **Summary table**
   Used to store pre-aggregated results.

```
CREATE TABLE TotalCheckoutsByBook AS
SELECT
   book_id,
   COUNT(checkout_id) AS TotalCheckouts
FROM
   checkoutsfact
GROUP BY
   book_id;
```

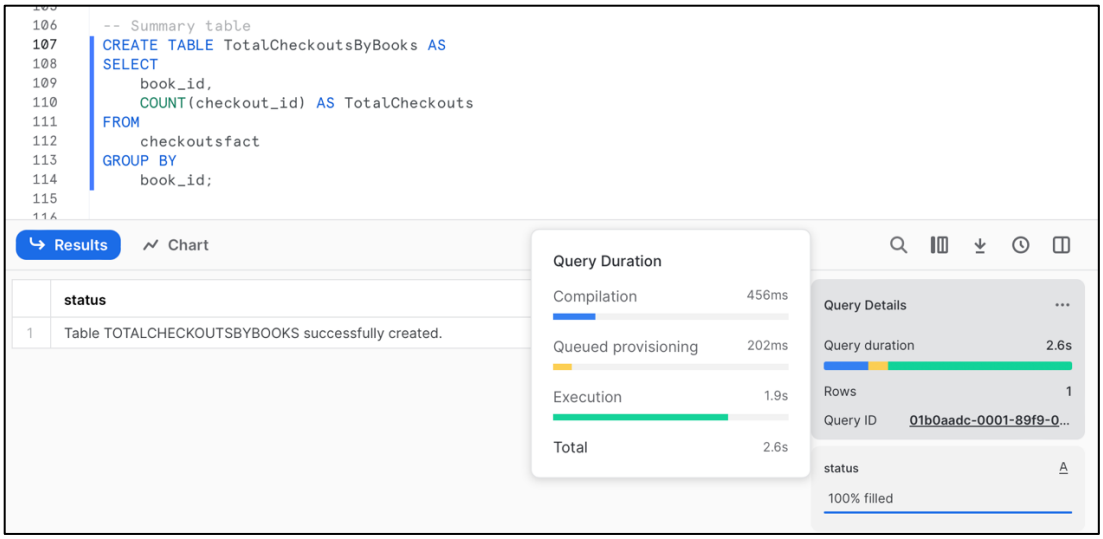

6. **Analysis**
   a. Identifying the most popular books based on the number of checkouts.

```
SELECT
   b.book_id,
   b.book_title,
```

```
    b.book_author,
    COUNT(c.checkout_id) AS checkout_count
FROM
    bookdim b
JOIN
    checkoutsfact c ON b.book_id = c.book_id
GROUP BY
    b.book_id, b.book_title, b.book_author
ORDER BY
    checkout_count DESC
LIMIT 10;
```

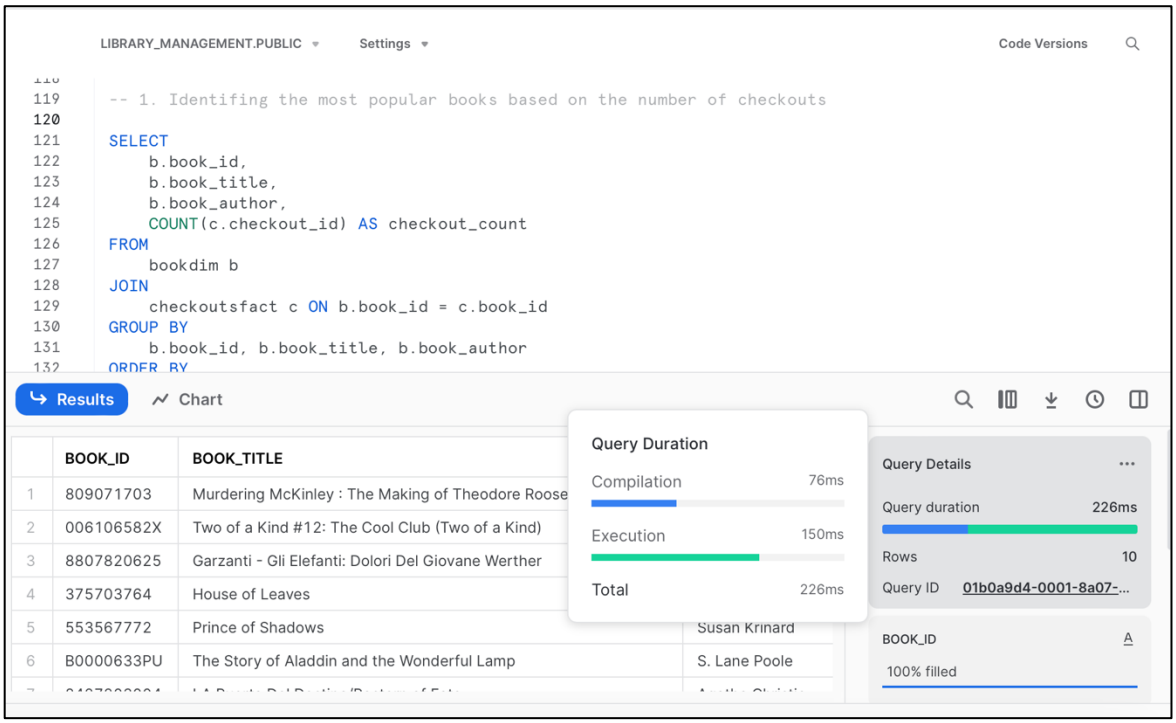

b. Calculated overdue fines and analyzed fine payment trends.

```
SELECT
    b.book_id,
    b.book_title,
    c.checkout_id,
    c.checkoutdate,
    c.returndate,
    f.paymentdate,
    f.amount
FROM
    checkoutsfact c
JOIN
    finepaymentsfact f ON c.checkout_id = f.checkoutid
```

```
JOIN
    bookdim b ON c.book_id = b.book_id
WHERE
    f.paymentdate > c.returndate;
```

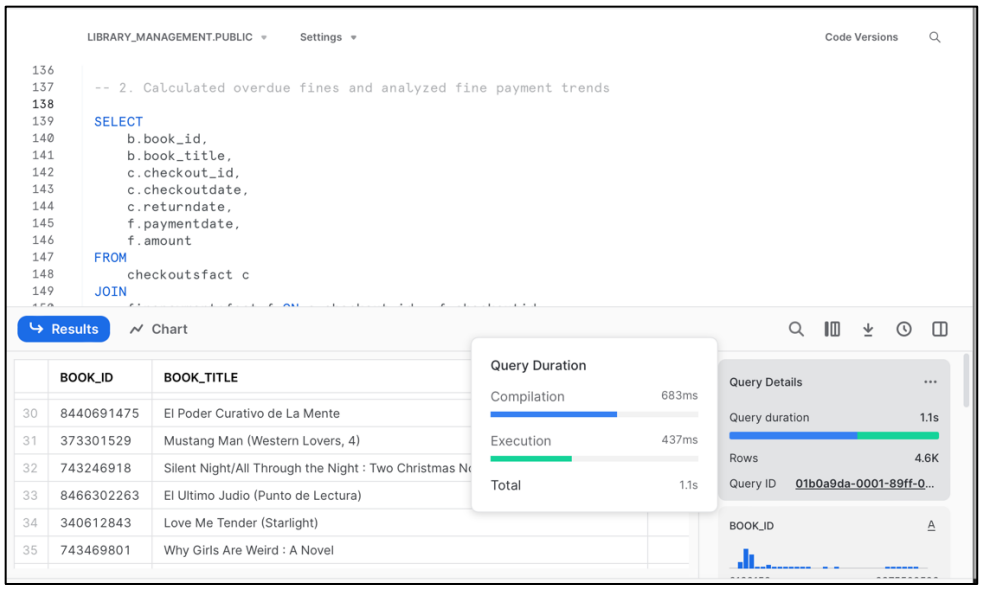

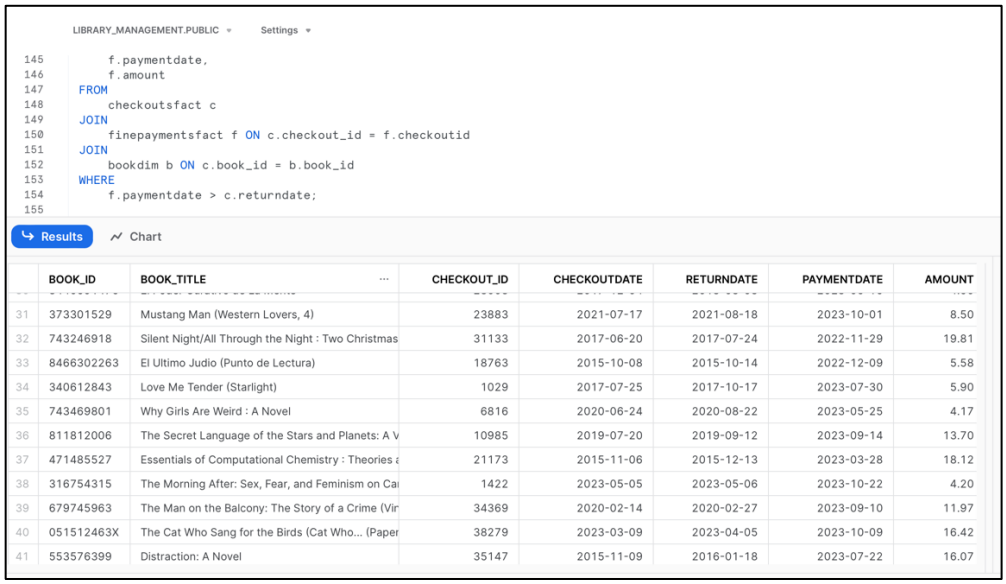

c. Number of times a User has borrowed books.
This query will provide a list of borrower IDs, borrower names, fact borrower IDs, and the count of records where there is a match between the BORROWER_ID values in the BORROWERDIM1 dimension table and the CHECKOUTSFACT fact table. It will be ordered by the borrower names.

```
SELECT
```

```
    B.BORROWER_ID AS DIM_BORROWER_ID,
    B.BORROWER_NAME,
    C.BORROWER_ID AS FACT_BORROWER_ID,
    COUNT(*) AS COUNT_OF_RECORDS
FROM
    BORROWERDIM1 B
LEFT JOIN
    CHECKOUTSFACT C ON B.BORROWER_ID = C.BORROWER_ID
GROUP BY
    B.BORROWER_ID, B.BORROWER_NAME, C.BORROWER_ID;
```

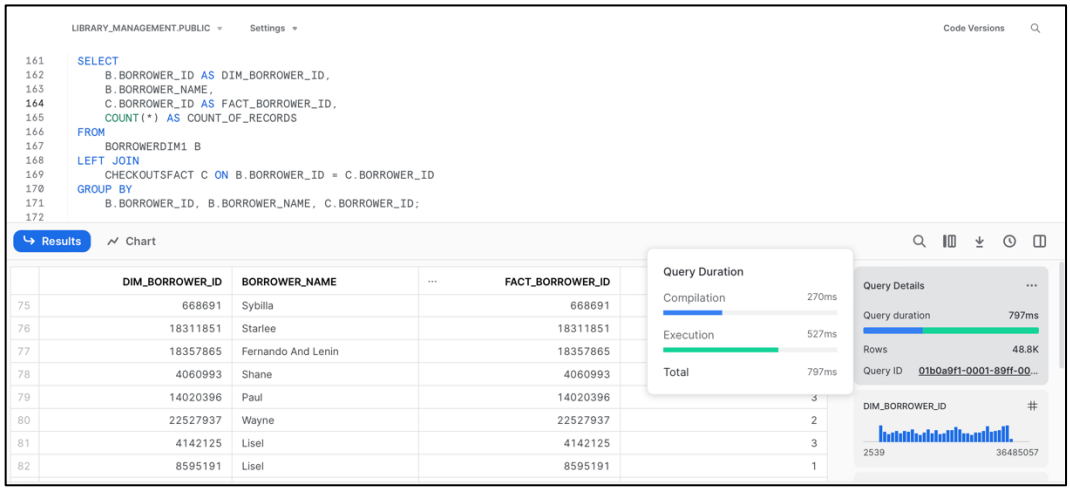

d. Views

This view aggregates data from the BORROWERDIM1, CHECKOUTSFACT, and FINEPAYMENTSFACT tables to provide borrower statistics, calculating the number of checkouts and total fine amounts for each borrower.

```
CREATE VIEW BorrowerStatisticsView AS
SELECT
    B.BORROWER_ID,
    B.BORROWER_NAME,
    COUNT(C.CHECKOUT_ID) AS NUM_CHECKOUTS,
    COALESCE(SUM(F.AMOUNT), 0) AS TOTAL_FINE_AMOUNT
FROM
    BORROWERDIM1 B
LEFT JOIN
    CHECKOUTSFACT C ON B.BORROWER_ID = C.BORROWER_ID
LEFT JOIN
    FINEPAYMENTSFACT F ON C.CHECKOUT_ID = F.CHECKOUTID
GROUP BY
    B.BORROWER_ID, B.BORROWER_NAME;
```

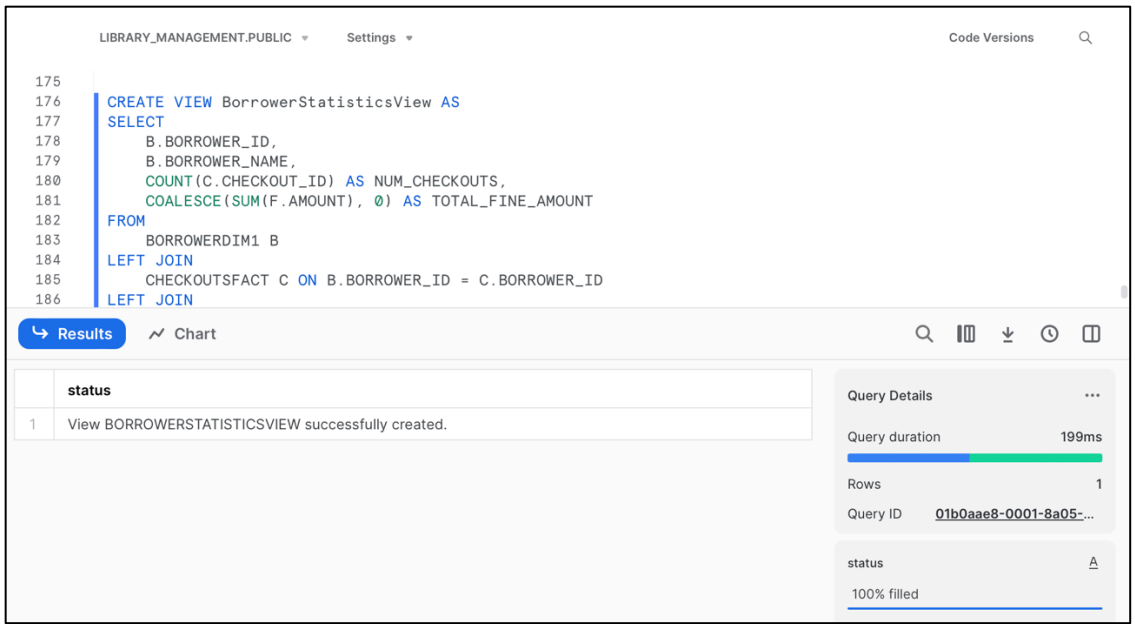

SELECT * from BorrowerStatisticsView;

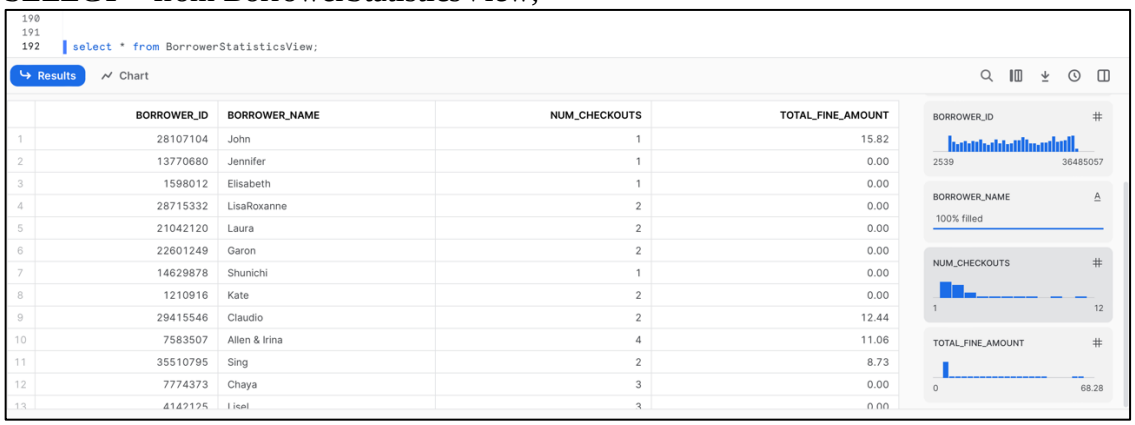

Dashboard: Created a Dashboard from Existing Worksheets

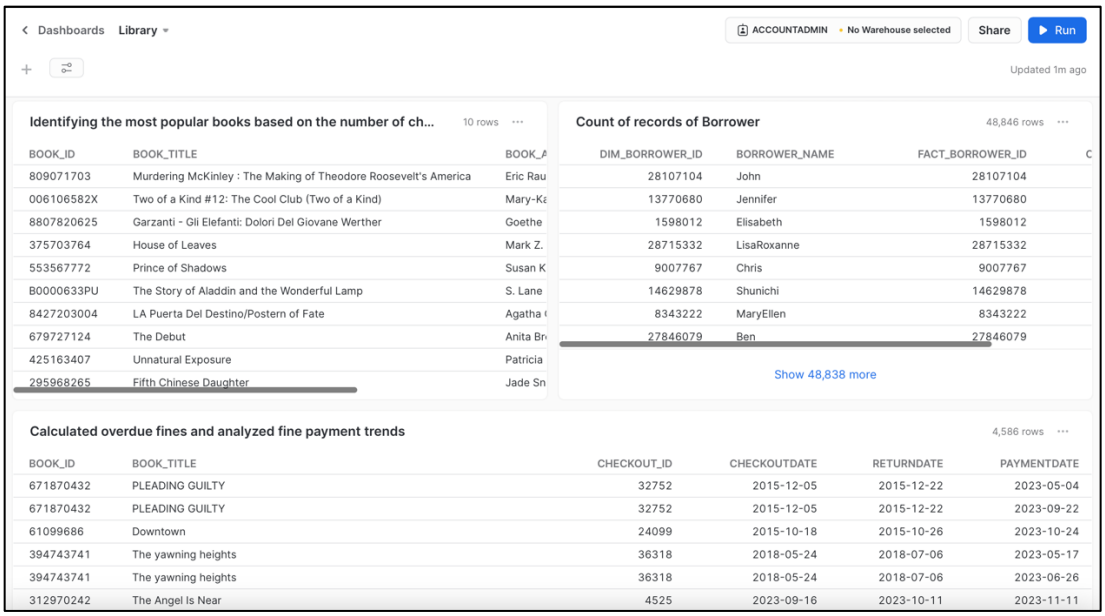

Snowflake's automatic scaling capability dynamically adapted resources in response to varying workload demands, guaranteeing peak performance and efficiency. This led to accelerated query execution and a more streamlined data processing experience.

The robust performance, security measures, and collaborative features of the platform established a strong basis for thorough data analysis and visualization, effectively accomplishing the project's goals.

Takeaways from this mini- project:

1. **Continuous Loading Using Snowpipe**
   It loads small volumes of data and periodically make them available for analysis. The process swiftly loads data within minutes of adding files to a designated stage, guaranteeing users access to the most recent results as soon as the raw data becomes accessible.

2. **Materialized views** are crafted to enhance query efficiency for workloads characterized by recurring and typical query patterns.

3. **Interacting with Secure Views**
   The view definition of a secure view is restricted to authorized users who have been granted the corresponding role ownership. Secure views safeguard users from potential exposure to filtered table rows, yet careful construction is crucial to avoid inadvertent data exposure.

4. Snowflake functions on a **multi-cloud framework**, enabling users to utilize cloud services such as AWS, Azure, and Google Cloud.

5. The system **segregates storage and computing resources**, offering flexibility, scalability, and cost efficiency in the administration of data and analytics workloads.

6. Snowflake supports **unstructured data** as well as natively supports **semi-structured data** formats like JSON.

7. **Multi-cluster Warehouse** enable you to scale compute resources to manage your user and query concurrency needs as they change, such as during peak and off hours.

8. Snowflake supports both **scaling up** by resizing a warehouse and **scaling out** by adding clusters to the warehouse. Snowflake also provides **automatic scaling**, adjusting resources dynamically according to workload requirements to guarantee optimal performance and efficiency.

9. Snowflake supports **temporary** (for storing non-permanent data) and **transient tables** until explicitly dropped.

10. **Search access path** monitors the potential presence of table column values within individual micro-partitions, allowing for the possibility of skipping certain micro-partitions during the table scanning process.

Blog link:
https://medium.com/@vaidehi.patel_164/library-management-system-using-snowflake-data-platform-e73b94fe390b